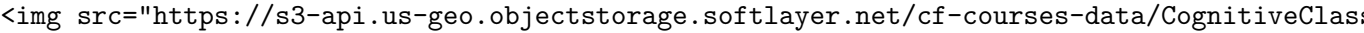


# exploratory-data-analysis

June 21, 2020

[<a href="https://cocl.us/corsera\\_da0101en\\_notebook\\_top">](https://cocl.us/corsera_da0101en_notebook_top)  
The image shows a browser window with a URL bar containing "https://s3-api.us-geo.objectstorage.softlayer.net/cf-courses-data/CognitiveClass/DA0101EN/automobileEDA.csv".  
</a>

Data Analysis with Python

Exploratory Data Analysis

Welcome!

In this section, we will explore several methods to see if certain characteristics or features can be used to predict car price.

Table of content

Import Data from Module

Analyzing Individual Feature Patterns using Visualization

Descriptive Statistical Analysis

Basics of Grouping

Correlation and Causation

ANOVA

Estimated Time Needed: 30 min

What are the main characteristics which have the most impact on the car price?

1. Import Data from Module 2

Setup

Import libraries

```
[1]: import pandas as pd
import numpy as np
```

load data and store in dataframe df:

This dataset was hosted on IBM Cloud object click [HERE](#) for free storage

```
[2]: path='https://s3-api.us-geo.objectstorage.softlayer.net/cf-courses-data/
      ↪CognitiveClass/DA0101EN/automobileEDA.csv'
df = pd.read_csv(path)
```

```
df.head()
```

```
[2]:   symboling  normalized-losses      make aspiration num-of-doors \
0         3           122 alfa-romero      std         two
1         3           122 alfa-romero      std         two
2         1           122 alfa-romero      std         two
3         2           164      audi      std         four
4         2           164      audi      std         four

      body-style drive-wheels engine-location  wheel-base  length  ... \
0  convertible      rwd      front      88.6  0.811148  ...
1  convertible      rwd      front      88.6  0.811148  ...
2   hatchback      rwd      front      94.5  0.822681  ...
3      sedan      fwd      front      99.8  0.848630  ...
4      sedan      4wd      front      99.4  0.848630  ...

      compression-ratio  horsepower  peak-rpm  city-mpg  highway-mpg  price \
0              9.0        111.0    5000.0      21          27  13495.0
1              9.0        111.0    5000.0      21          27  16500.0
2              9.0        154.0    5000.0      19          26  16500.0
3             10.0        102.0    5500.0      24          30  13950.0
4              8.0        115.0    5500.0      18          22  17450.0

      city-L/100km  horsepower-binned  diesel  gas
0    11.190476      Medium          0      1
1    11.190476      Medium          0      1
2    12.368421      Medium          0      1
3     9.791667      Medium          0      1
4    13.055556      Medium          0      1
```

[5 rows x 29 columns]

## 2. Analyzing Individual Feature Patterns using Visualization

To install seaborn we use the pip which is the python package manager.

```
[3]: %%capture
      ! pip install seaborn
```

Import visualization packages “Matplotlib” and “Seaborn”, don’t forget about “%matplotlib inline” to plot in a Jupyter notebook.

```
[4]: import matplotlib.pyplot as plt
      import seaborn as sns
      %matplotlib inline
```

How to choose the right visualization method?

When visualizing individual variables, it is important to first understand what type of variable you

are dealing with. This will help us find the right visualization method for that variable.

```
[5]: # list the data types for each column
      print(df.dtypes)
```

```
symboling          int64
normalized-losses  int64
make              object
aspiration         object
num-of-doors       object
body-style         object
drive-wheels       object
engine-location    object
wheel-base        float64
length            float64
width             float64
height            float64
curb-weight        int64
engine-type        object
num-of-cylinders   object
engine-size        int64
fuel-system        object
bore              float64
stroke            float64
compression-ratio  float64
horsepower         float64
peak-rpm          float64
city-mpg           int64
highway-mpg        int64
price             float64
city-L/100km       float64
horsepower-binned  object
diesel            int64
gas               int64
dtype: object
```

Question #1:

What is the data type of the column “peak-rpm”?

[Double-click here for the solution.](#)

for example, we can calculate the correlation between variables of type “int64” or “float64” using the method “corr”:

```
[6]: df.corr()
```

```
[6]:
```

	symboling	normalized-losses	wheel-base	length	\
symboling	1.000000	0.466264	-0.535987	-0.365404	
normalized-losses	0.466264	1.000000	-0.056661	0.019424	

wheel-base	-0.535987	-0.056661	1.000000	0.876024
length	-0.365404	0.019424	0.876024	1.000000
width	-0.242423	0.086802	0.814507	0.857170
height	-0.550160	-0.373737	0.590742	0.492063
curb-weight	-0.233118	0.099404	0.782097	0.880665
engine-size	-0.110581	0.112360	0.572027	0.685025
bore	-0.140019	-0.029862	0.493244	0.608971
stroke	-0.008245	0.055563	0.158502	0.124139
compression-ratio	-0.182196	-0.114713	0.250313	0.159733
horsepower	0.075819	0.217299	0.371147	0.579821
peak-rpm	0.279740	0.239543	-0.360305	-0.285970
city-mpg	-0.035527	-0.225016	-0.470606	-0.665192
highway-mpg	0.036233	-0.181877	-0.543304	-0.698142
price	-0.082391	0.133999	0.584642	0.690628
city-L/100km	0.066171	0.238567	0.476153	0.657373
diesel	-0.196735	-0.101546	0.307237	0.211187
gas	0.196735	0.101546	-0.307237	-0.211187

	width	height	curb-weight	engine-size	bore \
symboling	-0.242423	-0.550160	-0.233118	-0.110581	-0.140019
normalized-losses	0.086802	-0.373737	0.099404	0.112360	-0.029862
wheel-base	0.814507	0.590742	0.782097	0.572027	0.493244
length	0.857170	0.492063	0.880665	0.685025	0.608971
width	1.000000	0.306002	0.866201	0.729436	0.544885
height	0.306002	1.000000	0.307581	0.074694	0.180449
curb-weight	0.866201	0.307581	1.000000	0.849072	0.644060
engine-size	0.729436	0.074694	0.849072	1.000000	0.572609
bore	0.544885	0.180449	0.644060	0.572609	1.000000
stroke	0.188829	-0.062704	0.167562	0.209523	-0.055390
compression-ratio	0.189867	0.259737	0.156433	0.028889	0.001263
horsepower	0.615077	-0.087027	0.757976	0.822676	0.566936
peak-rpm	-0.245800	-0.309974	-0.279361	-0.256733	-0.267392
city-mpg	-0.633531	-0.049800	-0.749543	-0.650546	-0.582027
highway-mpg	-0.680635	-0.104812	-0.794889	-0.679571	-0.591309
price	0.751265	0.135486	0.834415	0.872335	0.543155
city-L/100km	0.673363	0.003811	0.785353	0.745059	0.554610
diesel	0.244356	0.281578	0.221046	0.070779	0.054458
gas	-0.244356	-0.281578	-0.221046	-0.070779	-0.054458

	stroke	compression-ratio	horsepower	peak-rpm \
symboling	-0.008245	-0.182196	0.075819	0.279740
normalized-losses	0.055563	-0.114713	0.217299	0.239543
wheel-base	0.158502	0.250313	0.371147	-0.360305
length	0.124139	0.159733	0.579821	-0.285970
width	0.188829	0.189867	0.615077	-0.245800
height	-0.062704	0.259737	-0.087027	-0.309974
curb-weight	0.167562	0.156433	0.757976	-0.279361

engine-size	0.209523	0.028889	0.822676	-0.256733
bore	-0.055390	0.001263	0.566936	-0.267392
stroke	1.000000	0.187923	0.098462	-0.065713
compression-ratio	0.187923	1.000000	-0.214514	-0.435780
horsepower	0.098462	-0.214514	1.000000	0.107885
peak-rpm	-0.065713	-0.435780	0.107885	1.000000
city-mpg	-0.034696	0.331425	-0.822214	-0.115413
highway-mpg	-0.035201	0.268465	-0.804575	-0.058598
price	0.082310	0.071107	0.809575	-0.101616
city-L/100km	0.037300	-0.299372	0.889488	0.115830
diesel	0.241303	0.985231	-0.169053	-0.475812
gas	-0.241303	-0.985231	0.169053	0.475812

	city-mpg	highway-mpg	price	city-L/100km	diesel \
symboling	-0.035527	0.036233	-0.082391	0.066171	-0.196735
normalized-losses	-0.225016	-0.181877	0.133999	0.238567	-0.101546
wheel-base	-0.470606	-0.543304	0.584642	0.476153	0.307237
length	-0.665192	-0.698142	0.690628	0.657373	0.211187
width	-0.633531	-0.680635	0.751265	0.673363	0.244356
height	-0.049800	-0.104812	0.135486	0.003811	0.281578
curb-weight	-0.749543	-0.794889	0.834415	0.785353	0.221046
engine-size	-0.650546	-0.679571	0.872335	0.745059	0.070779
bore	-0.582027	-0.591309	0.543155	0.554610	0.054458
stroke	-0.034696	-0.035201	0.082310	0.037300	0.241303
compression-ratio	0.331425	0.268465	0.071107	-0.299372	0.985231
horsepower	-0.822214	-0.804575	0.809575	0.889488	-0.169053
peak-rpm	-0.115413	-0.058598	-0.101616	0.115830	-0.475812
city-mpg	1.000000	0.972044	-0.686571	-0.949713	0.265676
highway-mpg	0.972044	1.000000	-0.704692	-0.930028	0.198690
price	-0.686571	-0.704692	1.000000	0.789898	0.110326
city-L/100km	-0.949713	-0.930028	0.789898	1.000000	-0.241282
diesel	0.265676	0.198690	0.110326	-0.241282	1.000000
gas	-0.265676	-0.198690	-0.110326	0.241282	-1.000000

	gas
symboling	0.196735
normalized-losses	0.101546
wheel-base	-0.307237
length	-0.211187
width	-0.244356
height	-0.281578
curb-weight	-0.221046
engine-size	-0.070779
bore	-0.054458
stroke	-0.241303
compression-ratio	-0.985231
horsepower	0.169053

peak-rpm	0.475812
city-mpg	-0.265676
highway-mpg	-0.198690
price	-0.110326
city-L/100km	0.241282
diesel	-1.000000
gas	1.000000

The diagonal elements are always one; we will study correlation more precisely Pearson correlation in-depth at the end of the notebook.

Question #2:

Find the correlation between the following columns: bore, stroke, compression-ratio, and horsepower.

Hint: if you would like to select those columns use the following syntax: `df[['bore','stroke','compression-ratio','horsepower']]`

```
[8]: # Write your code below and press Shift+Enter to execute
df[['bore','stroke','compression-ratio','horsepower']].corr()
```

```
[8]:          bore      stroke  compression-ratio  horsepower
bore          1.000000 -0.055390           0.001263    0.566936
stroke        -0.055390  1.000000           0.187923    0.098462
compression-ratio  0.001263  0.187923           1.000000   -0.214514
horsepower        0.566936  0.098462          -0.214514    1.000000
```

Double-click here for the solution.

Continuous numerical variables:

Continuous numerical variables are variables that may contain any value within some range. Continuous numerical variables can have the type “int64” or “float64”. A great way to visualize these variables is by using scatterplots with fitted lines.

In order to start understanding the (linear) relationship between an individual variable and the price. We can do this by using “regplot”, which plots the scatterplot plus the fitted regression line for the data.


Let’s see several examples of different linear relationships:

Positive linear relationship

Let’s find the scatterplot of “engine-size” and “price”

```
[9]: # Engine size as potential predictor variable of price
sns.regplot(x="engine-size", y="price", data=df)
plt.ylim(0,)
```

```
[9]: (0, 56118.75158672933)
```



As the engine-size goes up, the price goes up: this indicates a positive direct correlation between these two variables. Engine size seems like a pretty good predictor of price since the regression line is almost a perfect diagonal line.

We can examine the correlation between 'engine-size' and 'price' and see it's approximately 0.87


```
[10]: df[["engine-size", "price"]].corr()
```

```
[10]:      engine-size    price
engine-size    1.000000  0.872335
price          0.872335  1.000000
```

Highway mpg is a potential predictor variable of price

```
[11]: sns.regplot(x="highway-mpg", y="price", data=df)
```

```
[11]: <matplotlib.axes._subplots.AxesSubplot at 0x7f82103bc2b0>
```



As the highway-mpg goes up, the price goes down: this indicates an inverse/negative relationship between these two variables. Highway mpg could potentially be a predictor of price.

We can examine the correlation between 'highway-mpg' and 'price' and see it's approximately -0.704

```
[12]: df[['highway-mpg', 'price']].corr()
```

```
[12]:      highway-mpg    price
highway-mpg    1.000000 -0.704692
price          -0.704692  1.000000
```


Weak Linear Relationship

Let's see if "Peak-rpm" as a predictor variable of "price".

```
[13]: sns.regplot(x="peak-rpm", y="price", data=df)
```

```
[13]: <matplotlib.axes._subplots.AxesSubplot at 0x7f820c1ab940>
```





Peak rpm does not seem like a good predictor of the price at all since the regression line is close to horizontal. Also, the data points are very scattered and far from the fitted line, showing lots of variability. Therefore it's not a reliable variable.

We can examine the correlation between 'peak-rpm' and 'price' and see it's approximately -0.101616

```
[14]: df[['peak-rpm', 'price']].corr()
```

```
[14]:      peak-rpm    price
peak-rpm  1.000000 -0.101616
price    -0.101616  1.000000
```

Question 3 a):

Find the correlation between x="stroke", y="price".

Hint: if you would like to select those columns use the following syntax: `df[["stroke", "price"]]`

```
[15]: # Write your code below and press Shift+Enter to execute
df[['stroke', 'price']].corr()
```

```
[15]:      stroke    price
stroke  1.00000  0.08231
price   0.08231  1.00000
```


Double-click here for the solution.

Question 3 b):

Given the correlation results between “price” and “stroke” do you expect a linear relationship? Verify your results using the function “regplot()”.

```
[16]: # Write your code below and press Shift+Enter to execute
sns.regplot(x="stroke",y ="price", data = df)
```

```
[16]: <matplotlib.axes._subplots.AxesSubplot at 0x7f820c190f98>
```



[Double-click here for the solution.](#)

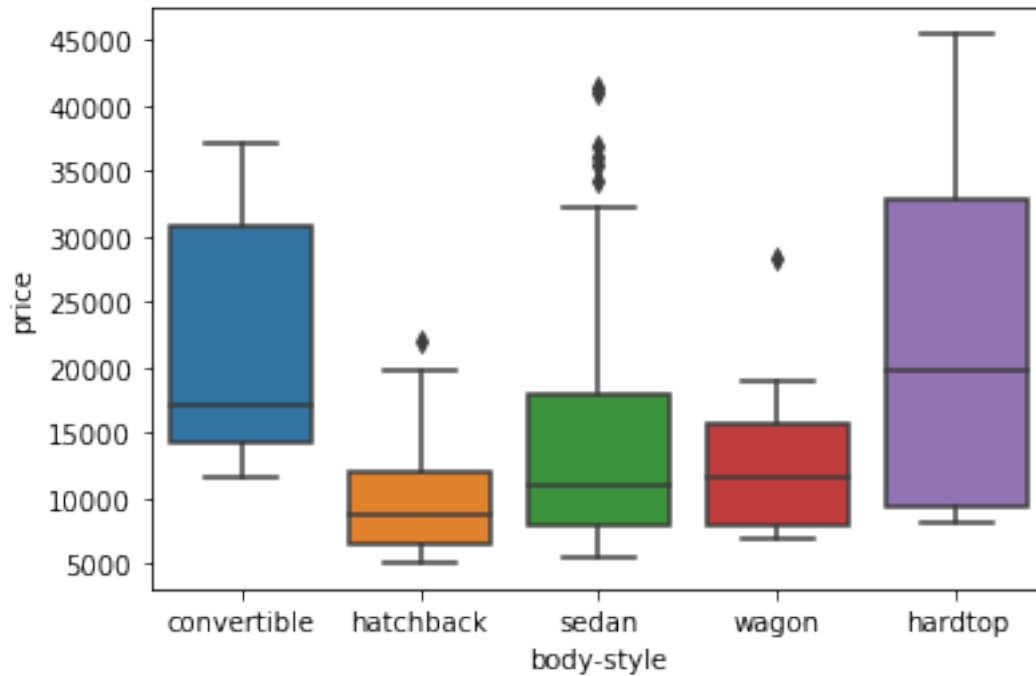
Categorical variables

These are variables that describe a ‘characteristic’ of a data unit, and are selected from a small group of categories. The categorical variables can have the type “object” or “int64”. A good way to visualize categorical variables is by using boxplots.

Let’s look at the relationship between “body-style” and “price”.

```
[17]: sns.boxplot(x="body-style", y="price", data=df)
```

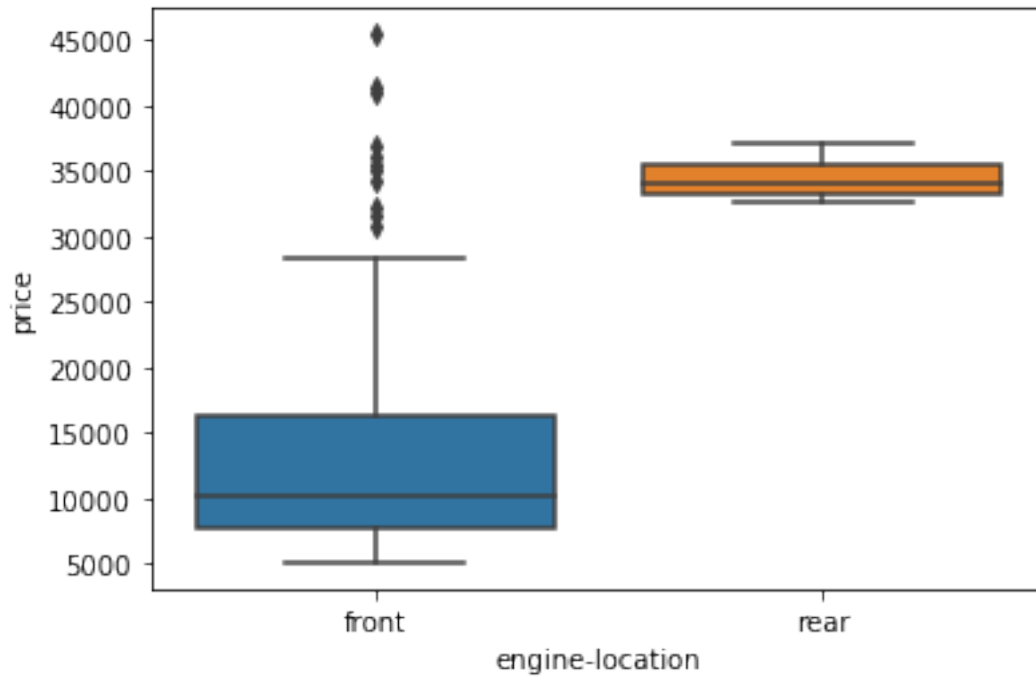
```
[17]: <matplotlib.axes._subplots.AxesSubplot at 0x7f82103ee940>
```



We see that the distributions of price between the different body-style categories have a significant overlap, and so body-style would not be a good predictor of price. Let's examine engine "engine-location" and "price":

```
[18]: sns.boxplot(x="engine-location", y="price", data=df)
```

```
[18]: <matplotlib.axes._subplots.AxesSubplot at 0x7f820c05f748>
```

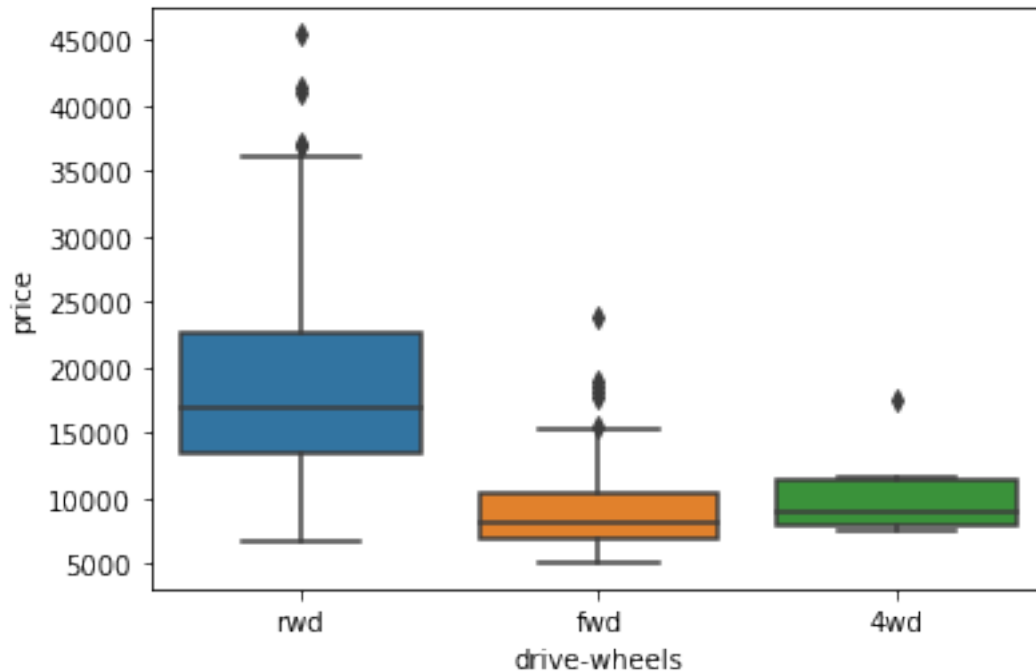


Here we see that the distribution of price between these two engine-location categories, front and rear, are distinct enough to take engine-location as a potential good predictor of price.

Let's examine "drive-wheels" and "price".

```
[19]: # drive-wheels
sns.boxplot(x="drive-wheels", y="price", data=df)
```

```
[19]: <matplotlib.axes._subplots.AxesSubplot at 0x7f82047b9f28>
```



Here we see that the distribution of price between the different drive-wheels categories differs; as such drive-wheels could potentially be a predictor of price.

### 3. Descriptive Statistical Analysis

Let's first take a look at the variables by utilizing a description method.

The describe function automatically computes basic statistics for all continuous variables. Any NaN values are automatically skipped in these statistics.

This will show:

the count of that variable

the mean

the standard deviation (std)

the minimum value

the IQR (Interquartile Range: 25%, 50% and 75%)

the maximum value

We can apply the method “describe” as follows:

```
[20]: df.describe()
```

```
[20]:
```

	symboling	normalized-losses	wheel-base	length	width	\
count	201.000000	201.00000	201.000000	201.000000	201.000000	

mean	0.840796	122.00000	98.797015	0.837102	0.915126
std	1.254802	31.99625	6.066366	0.059213	0.029187
min	-2.000000	65.00000	86.600000	0.678039	0.837500
25%	0.000000	101.00000	94.500000	0.801538	0.890278
50%	1.000000	122.00000	97.000000	0.832292	0.909722
75%	2.000000	137.00000	102.400000	0.881788	0.925000
max	3.000000	256.00000	120.900000	1.000000	1.000000

	height	curb-weight	engine-size	bore	stroke \
count	201.000000	201.000000	201.000000	201.000000	197.000000
mean	53.766667	2555.666667	126.875622	3.330692	3.256904
std	2.447822	517.296727	41.546834	0.268072	0.319256
min	47.800000	1488.000000	61.000000	2.540000	2.070000
25%	52.000000	2169.000000	98.000000	3.150000	3.110000
50%	54.100000	2414.000000	120.000000	3.310000	3.290000
75%	55.500000	2926.000000	141.000000	3.580000	3.410000
max	59.800000	4066.000000	326.000000	3.940000	4.170000

	compression-ratio	horsepower	peak-rpm	city-mpg	highway-mpg \
count	201.000000	201.000000	201.000000	201.000000	201.000000
mean	10.164279	103.405534	5117.665368	25.179104	30.686567
std	4.004965	37.365700	478.113805	6.423220	6.815150
min	7.000000	48.000000	4150.000000	13.000000	16.000000
25%	8.600000	70.000000	4800.000000	19.000000	25.000000
50%	9.000000	95.000000	5125.369458	24.000000	30.000000
75%	9.400000	116.000000	5500.000000	30.000000	34.000000
max	23.000000	262.000000	6600.000000	49.000000	54.000000

	price	city-L/100km	diesel	gas
count	201.000000	201.000000	201.000000	201.000000
mean	13207.129353	9.944145	0.099502	0.900498
std	7947.066342	2.534599	0.300083	0.300083
min	5118.000000	4.795918	0.000000	0.000000
25%	7775.000000	7.833333	0.000000	1.000000
50%	10295.000000	9.791667	0.000000	1.000000
75%	16500.000000	12.368421	0.000000	1.000000
max	45400.000000	18.076923	1.000000	1.000000

The default setting of “describe” skips variables of type object. We can apply the method “describe” on the variables of type ‘object’ as follows:

```
[21]: df.describe(include=['object'])
```

```
[21]:
```

	make	aspiration	num-of-doors	body-style	drive-wheels \
count	201	201	201	201	201
unique	22	2	2	5	3
top	toyota	std	four	sedan	fwd

freq	32	165	115	94	118
------	----	-----	-----	----	-----

	engine-location	engine-type	num-of-cylinders	fuel-system	\
count	201	201	201	201	
unique	2	6	7	8	
top	front	ohc	four	mpfi	
freq	198	145	157	92	

	horsepower-binned
count	200
unique	3
top	Low
freq	115

### Value Counts

Value-counts is a good way of understanding how many units of each characteristic/variable we have. We can apply the “value\_counts” method on the column ‘drive-wheels’. Don’t forget the method “value\_counts” only works on Pandas series, not Pandas Dataframes. As a result, we only include one bracket “df[‘drive-wheels’]” not two brackets “df[[‘drive-wheels’]]”.

```
[22]: df['drive-wheels'].value_counts()
```

```
[22]: fwd      118
      rwd       75
      4wd        8
      Name: drive-wheels, dtype: int64
```

We can convert the series to a Dataframe as follows :

```
[23]: df['drive-wheels'].value_counts().to_frame()
```

```
[23]:      drive-wheels
      fwd      118
      rwd      75
      4wd       8
```

Let’s repeat the above steps but save the results to the dataframe “drive\_wheels\_counts” and rename the column ‘drive-wheels’ to ‘value\_counts’.

```
[24]: drive_wheels_counts = df['drive-wheels'].value_counts().to_frame()
      drive_wheels_counts.rename(columns={'drive-wheels': 'value_counts'},
      inplace=True)
      drive_wheels_counts
```

```
[24]:      value_counts
      fwd      118
      rwd      75
      4wd       8
```

Now let's rename the index to 'drive-wheels':

```
[25]: drive_wheels_counts.index.name = 'drive-wheels'
      drive_wheels_counts
```

```
[25]:          value_counts
drive-wheels
fwd          118
rwd           75
4wd           8
```

We can repeat the above process for the variable 'engine-location'.

```
[26]: # engine-location as variable
      engine_loc_counts = df['engine-location'].value_counts().to_frame()
      engine_loc_counts.rename(columns={'engine-location': 'value_counts'},
      →inplace=True)
      engine_loc_counts.index.name = 'engine-location'
      engine_loc_counts.head(10)
```

```
[26]:          value_counts
engine-location
front          198
rear            3
```

Examining the value counts of the engine location would not be a good predictor variable for the price. This is because we only have three cars with a rear engine and 198 with an engine in the front, this result is skewed. Thus, we are not able to draw any conclusions about the engine location.

#### 4. Basics of Grouping

The “groupby” method groups data by different categories. The data is grouped based on one or several variables and analysis is performed on the individual groups.

For example, let's group by the variable “drive-wheels”. We see that there are 3 different categories of drive wheels.

```
[27]: df['drive-wheels'].unique()
```

```
[27]: array(['rwd', 'fwd', '4wd'], dtype=object)
```

If we want to know, on average, which type of drive wheel is most valuable, we can group “drive-wheels” and then average them.

We can select the columns 'drive-wheels', 'body-style' and 'price', then assign it to the variable “df\_group\_one”.

```
[29]: df_group_one = df[['drive-wheels', 'body-style', 'price']]
```

We can then calculate the average price for each of the different categories of data.



```
[30]: # grouping results
df_group_one = df_group_one.groupby(['drive-wheels'],as_index=False).mean()
df_group_one
```

```
[30]:   drive-wheels      price
0         4wd  10241.000000
1         fwd   9244.779661
2         rwd  19757.613333
```

From our data, it seems rear-wheel drive vehicles are, on average, the most expensive, while 4-wheel and front-wheel are approximately the same in price.

You can also group with multiple variables. For example, let's group by both 'drive-wheels' and 'body-style'. This groups the dataframe by the unique combinations 'drive-wheels' and 'body-style'. We can store the results in the variable 'grouped\_test1'.

```
[31]: # grouping results
df_gptest = df[['drive-wheels','body-style','price']]
grouped_test1 = df_gptest.groupby(['drive-wheels','body-style'],as_index=False).
    ↪mean()
grouped_test1
```

```
[31]:   drive-wheels  body-style      price
0         4wd    hatchback  7603.000000
1         4wd      sedan   12647.333333
2         4wd      wagon   9095.750000
3         fwd  convertible  11595.000000
4         fwd    hardtop   8249.000000
5         fwd    hatchback  8396.387755
6         fwd      sedan   9811.800000
7         fwd      wagon   9997.333333
8         rwd  convertible  23949.600000
9         rwd    hardtop   24202.714286
10        rwd    hatchback  14337.777778
11        rwd      sedan   21711.833333
12        rwd      wagon   16994.222222
```

This grouped data is much easier to visualize when it is made into a pivot table. A pivot table is like an Excel spreadsheet, with one variable along the column and another along the row. We can convert the dataframe to a pivot table using the method "pivot" to create a pivot table from the groups.

In this case, we will leave the drive-wheel variable as the rows of the table, and pivot body-style to become the columns of the table:

```
[32]: grouped_pivot = grouped_test1.pivot(index='drive-wheels',columns='body-style')
grouped_pivot
```

```
[32]:
```

	price			
body-style	convertible	hardtop	hatchback	sedan
drive-wheels				
4wd	NaN	NaN	7603.000000	12647.333333
fwd	11595.0	8249.000000	8396.387755	9811.800000
rwd	23949.6	24202.714286	14337.777778	21711.833333

body-style	wagon
drive-wheels	
4wd	9095.750000
fwd	9997.333333
rwd	16994.222222

Often, we won't have data for some of the pivot cells. We can fill these missing cells with the value 0, but any other value could potentially be used as well. It should be mentioned that missing data is quite a complex subject and is an entire course on its own.

```
[33]: grouped_pivot = grouped_pivot.fillna(0) #fill missing values with 0
grouped_pivot
```

```
[33]:
```

	price			
body-style	convertible	hardtop	hatchback	sedan
drive-wheels				
4wd	0.0	0.000000	7603.000000	12647.333333
fwd	11595.0	8249.000000	8396.387755	9811.800000
rwd	23949.6	24202.714286	14337.777778	21711.833333

body-style	wagon
drive-wheels	
4wd	9095.750000
fwd	9997.333333
rwd	16994.222222

Question 4:

Use the “groupby” function to find the average “price” of each car based on “body-style” ?

```
[34]: # Write your code below and press Shift+Enter to execute
group2 = df[['body-style', 'price']]
grouped_test_bodystyle = group2.groupby(['body-style'], as_index=False).mean()
grouped_test_bodystyle
```

```
[34]:
```

	body-style	price
0	convertible	21890.500000
1	hardtop	22208.500000
2	hatchback	9957.441176

```
3      sedan  14459.755319
4      wagon  12371.960000
```

Double-click here for the solution.


If you did not import “pyplot” let’s do it again.

```
[35]: import matplotlib.pyplot as plt
      %matplotlib inline
```

Variables: Drive Wheels and Body Style vs Price

Let’s use a heat map to visualize the relationship between Body Style vs Price.

```
[36]: #use the grouped results
      plt.pcolor(grouped_pivot, cmap='RdBu')
      plt.colorbar()
      plt.show()
```



The heatmap plots the target variable (price) proportional to colour with respect to the variables ‘drive-wheel’ and ‘body-style’ in the vertical and horizontal axis respectively. This allows us to visualize how the price is related to ‘drive-wheel’ and ‘body-style’.

The default labels convey no useful information to us. Let’s change that:

```
[40]: fig, ax = plt.subplots()
      im = ax.pcolor(grouped_pivot, cmap='RdBu')
```

```

#label names
row_labels = grouped_pivot.columns.levels[1]
col_labels = grouped_pivot.index


#move ticks and labels to the center
ax.set_xticks(np.arange(grouped_pivot.shape[1]) + 0.5, minor=False)
ax.set_yticks(np.arange(grouped_pivot.shape[0]) + 0.5, minor=False)

#insert labels
ax.set_xticklabels(row_labels, minor=False)
ax.set_yticklabels(col_labels, minor=False)

#rotate label if too long
plt.xticks(rotation=90)

fig.colorbar(im)
plt.show()

```



Visualization is very important in data science, and Python visualization packages provide great freedom. We will go more in-depth in a separate Python Visualizations course.

The main question we want to answer in this module, is “What are the main characteristics which have the most impact on the car price?”.

To get a better measure of the important characteristics, we look at the correlation of these variables with the car price, in other words: how is the car price dependent on this variable?

## 5. Correlation and Causation

Correlation: a measure of the extent of interdependence between variables.

Causation: the relationship between cause and effect between two variables.

It is important to know the difference between these two and that correlation does not imply causation. Determining correlation is much simpler than determining causation as causation may require independent experimentation.

### Pearson Correlation

The Pearson Correlation measures the linear dependence between two variables X and Y.

The resulting coefficient is a value between -1 and 1 inclusive, where:

1: Total positive linear correlation.

0: No linear correlation, the two variables most likely do not affect each other.

-1: Total negative linear correlation.

Pearson Correlation is the default method of the function “corr”. Like before we can calculate the Pearson Correlation of the of the ‘int64’ or ‘float64’ variables.

```
[ ]: df.corr()
```

sometimes we would like to know the significant of the correlation estimate.

P-value:

What is this P-value? The P-value is the probability value that the correlation between these two variables is statistically significant. Normally, we choose a significance level of 0.05, which means that we are 95% confident that the correlation between the variables is significant.

By convention, when the

p-value is  $< 0.001$ : we say there is strong evidence that the correlation is significant.

the p-value is  $< 0.05$ : there is moderate evidence that the correlation is significant.

the p-value is  $< 0.1$ : there is weak evidence that the correlation is significant.

the p-value is  $> 0.1$ : there is no evidence that the correlation is significant.

We can obtain this information using “stats” module in the “scipy” library.

```
[41]: from scipy import stats
```

Wheel-base vs Price

Let’s calculate the Pearson Correlation Coefficient and P-value of ‘wheel-base’ and ‘price’.

```
[42]: pearson_coef, p_value = stats.pearsonr(df['wheel-base'], df['price'])
```

```
print("The Pearson Correlation Coefficient is", pearson_coef, " with a P-value_\n→of P =", p_value)
```

The Pearson Correlation Coefficient is 0.584641822265508 with a P-value of P = 8.076488270733218e-20

Conclusion:

Since the p-value is  $< 0.001$ , the correlation between wheel-base and price is statistically significant, although the linear relationship isn't extremely strong (~0.585)

Horsepower vs Price

Let's calculate the Pearson Correlation Coefficient and P-value of 'horsepower' and 'price'.

```
[43]: pearson_coef, p_value = stats.pearsonr(df['horsepower'], df['price'])
print("The Pearson Correlation Coefficient is", pearson_coef, " with a P-value_\n→of P = ", p_value)
```

The Pearson Correlation Coefficient is 0.8095745670036559 with a P-value of P = 6.369057428260101e-48

Conclusion:

Since the p-value is  $< 0.001$ , the correlation between horsepower and price is statistically significant, and the linear relationship is quite strong (~0.809, close to 1)

Length vs Price

Let's calculate the Pearson Correlation Coefficient and P-value of 'length' and 'price'.

```
[44]: pearson_coef, p_value = stats.pearsonr(df['length'], df['price'])
print("The Pearson Correlation Coefficient is", pearson_coef, " with a P-value_\n→of P = ", p_value)
```

The Pearson Correlation Coefficient is 0.6906283804483638 with a P-value of P = 8.016477466159556e-30

Conclusion:

Since the p-value is  $< 0.001$ , the correlation between length and price is statistically significant, and the linear relationship is moderately strong (~0.691).

Width vs Price

Let's calculate the Pearson Correlation Coefficient and P-value of 'width' and 'price':

```
[45]: pearson_coef, p_value = stats.pearsonr(df['width'], df['price'])
print("The Pearson Correlation Coefficient is", pearson_coef, " with a P-value_\n→of P =", p_value )
```

The Pearson Correlation Coefficient is 0.7512653440522673 with a P-value of P = 9.200335510481646e-38

**Conclusion:** Since the p-value is  $< 0.001$ , the correlation between width and price is statistically significant, and the linear relationship is quite strong ( $\sim 0.751$ ).

### 0.0.1 Curb-weight vs Price

Let's calculate the Pearson Correlation Coefficient and P-value of 'curb-weight' and 'price':

```
[46]: pearson_coef, p_value = stats.pearsonr(df['curb-weight'], df['price'])
      print( "The Pearson Correlation Coefficient is", pearson_coef, " with a P-value_
      ↪of P = ", p_value)
```

The Pearson Correlation Coefficient is 0.8344145257702843 with a P-value of P = 2.189577238894065e-53

Conclusion:

Since the p-value is  $< 0.001$ , the correlation between curb-weight and price is statistically significant, and the linear relationship is quite strong ( $\sim 0.834$ ).

Engine-size vs Price

Let's calculate the Pearson Correlation Coefficient and P-value of 'engine-size' and 'price':

```
[47]: pearson_coef, p_value = stats.pearsonr(df['engine-size'], df['price'])
      print("The Pearson Correlation Coefficient is", pearson_coef, " with a P-value_
      ↪of P =", p_value)
```

The Pearson Correlation Coefficient is 0.8723351674455185 with a P-value of P = 9.265491622198389e-64

Conclusion:

Since the p-value is  $< 0.001$ , the correlation between engine-size and price is statistically significant, and the linear relationship is very strong ( $\sim 0.872$ ).

Bore vs Price

Let's calculate the Pearson Correlation Coefficient and P-value of 'bore' and 'price':

```
[48]: pearson_coef, p_value = stats.pearsonr(df['bore'], df['price'])
      print("The Pearson Correlation Coefficient is", pearson_coef, " with a P-value_
      ↪of P = ", p_value )
```

The Pearson Correlation Coefficient is 0.5431553832626602 with a P-value of P = 8.049189483935489e-17

Conclusion:

Since the p-value is  $< 0.001$ , the correlation between bore and price is statistically significant, but the linear relationship is only moderate ( $\sim 0.521$ ).

We can relate the process for each 'City-mpg' and 'Highway-mpg':

City-mpg vs Price

```
[49]: pearson_coef, p_value = stats.pearsonr(df['city-mpg'], df['price'])
print("The Pearson Correlation Coefficient is", pearson_coef, " with a P-value of P = ", p_value)
```

The Pearson Correlation Coefficient is -0.6865710067844678 with a P-value of P = 2.321132065567641e-29

Conclusion:

Since the p-value is  $< 0.001$ , the correlation between city-mpg and price is statistically significant, and the coefficient of  $\sim -0.687$  shows that the relationship is negative and moderately strong.

Highway-mpg vs Price

```
[50]: pearson_coef, p_value = stats.pearsonr(df['highway-mpg'], df['price'])
print( "The Pearson Correlation Coefficient is", pearson_coef, " with a P-value of P = ", p_value )
```

The Pearson Correlation Coefficient is -0.704692265058953 with a P-value of P = 1.7495471144476358e-31

**Conclusion:** Since the p-value is  $< 0.001$ , the correlation between highway-mpg and price is statistically significant, and the coefficient of  $\sim -0.705$  shows that the relationship is negative and moderately strong.

## 6. ANOVA

ANOVA: Analysis of Variance

The Analysis of Variance (ANOVA) is a statistical method used to test whether there are significant differences between the means of two or more groups. ANOVA returns two parameters:

F-test score: ANOVA assumes the means of all groups are the same, calculates how much the actual means deviate from the assumption, and reports it as the F-test score. A larger score means there is a larger difference between the means.

P-value: P-value tells how statistically significant is our calculated score value.

If our price variable is strongly correlated with the variable we are analyzing, expect ANOVA to return a sizeable F-test score and a small p-value.

Drive Wheels

Since ANOVA analyzes the difference between different groups of the same variable, the groupby function will come in handy. Because the ANOVA algorithm averages the data automatically, we do not need to take the average before hand.

Let's see if different types 'drive-wheels' impact 'price', we group the data.

Let's see if different types 'drive-wheels' impact 'price', we group the data.

```
[51]: grouped_test2=df_gptest[['drive-wheels', 'price']].groupby(['drive-wheels'])
grouped_test2.head(2)
```



```
[51]:      drive-wheels    price
      0          rwd  13495.0
      1          rwd  16500.0
      3          fwd  13950.0
      4          4wd  17450.0
      5          fwd  15250.0
     136          4wd   7603.0
```

```
[52]: df_gptest
```

```
[52]:      drive-wheels  body-style    price
      0          rwd  convertible  13495.0
      1          rwd  convertible  16500.0
      2          rwd   hatchback  16500.0
      3          fwd     sedan    13950.0
      4          4wd     sedan    17450.0
      ..          ...          ...      ...
     196          rwd     sedan    16845.0
     197          rwd     sedan    19045.0
     198          rwd     sedan    21485.0
     199          rwd     sedan    22470.0
     200          rwd     sedan    22625.0
```

```
[201 rows x 3 columns]
```

We can obtain the values of the method group using the method “get\_group”.

```
[53]: grouped_test2.get_group('4wd')['price']
```

```
[53]: 4          17450.0
     136          7603.0
     140          9233.0
     141         11259.0
     144          8013.0
     145         11694.0
     150          7898.0
     151          8778.0
     Name: price, dtype: float64
```

we can use the function ‘f\_oneway’ in the module ‘stats’ to obtain the F-test score and P-value.

```
[54]: # ANOVA
f_val, p_val = stats.f_oneway(grouped_test2.get_group('fwd')['price'],
    ↳ grouped_test2.get_group('rwd')['price'], grouped_test2.
    ↳ get_group('4wd')['price'])

print( "ANOVA results: F=", f_val, ", P =", p_val)
```

ANOVA results: F= 67.95406500780399 , P = 3.3945443577151245e-23

This is a great result, with a large F test score showing a strong correlation and a P value of almost 0 implying almost certain statistical significance. But does this mean all three tested groups are all this highly correlated?

**Separately: fwd and rwd**

```
[55]: f_val, p_val = stats.f_oneway(grouped_test2.get_group('fwd')['price'],  
    ↪grouped_test2.get_group('rwd')['price'])  
  
print( "ANOVA results: F=", f_val, ", P =", p_val )
```

ANOVA results: F= 130.5533160959111 , P = 2.2355306355677845e-23

Let's examine the other groups

**4wd and rwd**

```
[56]: f_val, p_val = stats.f_oneway(grouped_test2.get_group('4wd')['price'],  
    ↪grouped_test2.get_group('rwd')['price'])  
  
print( "ANOVA results: F=", f_val, ", P =", p_val)
```

ANOVA results: F= 8.580681368924756 , P = 0.004411492211225333

4wd and fwd

```
[57]: f_val, p_val = stats.f_oneway(grouped_test2.get_group('4wd')['price'],  
    ↪grouped_test2.get_group('fwd')['price'])  
  
print("ANOVA results: F=", f_val, ", P =", p_val)
```

ANOVA results: F= 0.665465750252303 , P = 0.41620116697845666

Conclusion: Important Variables

We now have a better idea of what our data looks like and which variables are important to take into account when predicting the car price. We have narrowed it down to the following variables:

Continuous numerical variables:

Length

Width

Curb-weight

Engine-size

Horsepower

City-mpg

Highway-mpg

Wheel-base

Bore

Categorical variables:

Drive-wheels

As we now move into building machine learning models to automate our analysis, feeding the model with variables that meaningfully affect our target variable will improve our model's prediction performance.

Thank you for completing this notebook

<p><a href="https://cocl.us/corsera\_da0101en\_notebook\_bottom"><img src="https://s3-api.us-geo.

About the Authors:

This notebook was written by Mahdi Noorian PhD, Joseph Santarcangelo, Bahare Talayian, Eric Xiao, Steven Dong, Parizad, Hima Vsudevan and Fiorella Wenver and Yi Yao.

Joseph Santarcangelo is a Data Scientist at IBM, and holds a PhD in Electrical Engineering. His research focused on using Machine Learning, Signal Processing, and Computer Vision to determine how videos impact human cognition. Joseph has been working for IBM since he completed his PhD.

Copyright © 2018 IBM Developer Skills Network. This notebook and its source code are released under the terms of the MIT License.